Eric Pitman Summer Workshop in Computational Science

3. Descriptive Statistics

Jeanette Sperhac

Descriptive Statistics

Explore a dataset:

- What's in the dataset?

- What does it mean?

- What if there's *a lot* of it?

Basic Statistical Functions in R

Wanted: measures of the center and the spread of our numeric data.

- mean()
- median()
- range()
- var() and sd()   # variance, standard deviation
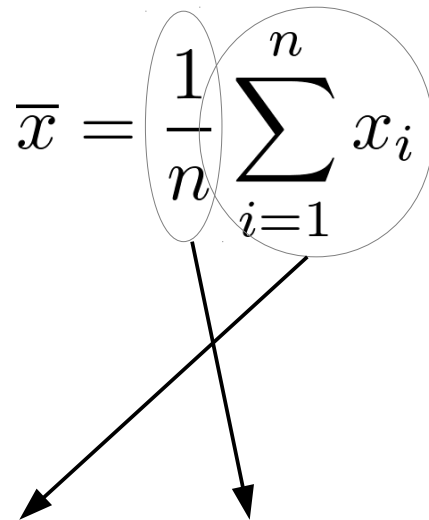- summary()   # combination of measures

mean()

A measure of the data's "most typical" value.

- Arithmetic mean == average
- Divide sum of values by number of values

$$\overline{x} = \frac{1}{n}\sum_{i=1}^{n} x_i$$

mean()

A measure of the data's "most typical" value.

$$\overline{x} = \frac{1}{n}\sum_{i=1}^{n} x_i$$

```
> f <- c(3, 2, 4, 1)
> mean(f)    # == sum(f)/length(f) == (3+2+4+1)/4
[1] 2.5
```

median()

A measure of the data's center value. To find it:

- Sort the contents of the data structure

- Compute the value at the center of the data:

    - For odd number of elements, take the center element's value.

    - For even number of elements, take mean around center.

median()

Odd number of values:

> h <- c(3, 1, 2)

> median(h)

[1] 2

median()

Even number of values: need to find mean()

f
1   3
2   2
3   4
4   1

sort →

f'
1   1
2   2
3   3
4   4

find center →

f'
1   1
2   2
3   3
4   4

← median(f)
= mean(c(2,3))
= 2.5

> f <- c(3, 2, 4, 1)

> median(f)

[1] 2.50

range():
min() and max()

range() reports the minimum and maximum values found in the data structure.

> f <- c(3, 2, 4, 1)

> range(f)   # reports min(f) and max(f)

   [1] 1   4

var() and sd()

- *Variance*: a measure of the spread of the values relative to their mean:

$$Var = s_n^2 = \frac{1}{n}\sum_{i=1}^{n}(y_i - \overline{y})^2$$ Sample variance

- *Standard deviation*: square root of the variance

$$s_n = \sqrt{Var}$$ Sample standard deviation

R's summary() Function

Provides several useful descriptive statistics about the data:

> g <- c(3, NA, 2, NA, 4, 1)

> summary(g)

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
1.00	1.75	2.50	2.50	3.25	4.00	2

*Quartiles*: Sort the data set and divide it up into quarters...

Quartiles

Quartiles are the *three points* that divide ordered data into four equal-sized groups:

- Q1 marks the boundary just above the lowest 25% of the data

- Q2 (the *median*) cuts the data set in half

- Q3 marks the boundary just below the highest 25% of data

Quartiles

Boxplot and probability distribution function of Normal N(0,1σ^2) population

Summary: Basic Statistical Functions

- Characterize the center and the spread of our numeric data.

- Comparing these measures can give us a good sense of our dataset.

Statistics and Missing Data

If NAs are present, specify na.rm=TRUE to call:

- mean()

- median()

- range()

- sum()

- ...and some other functions

R disregards NAs, then proceeds with the calculation.

Diamonds Data

50,000 diamonds, for example:

	carat	cut	color	clarity	depth	table	price	x	y	z
1	0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
2	0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
3	0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31

What can we learn about these data?

Diamonds Data

In the RStudio source editor:

<workshop>/examples/diamonds.csv

To load from the command line:

- Set the Working Directory

- Load command:

```
> d=read.csv("diamonds.csv")
```

Diamonds Data summary()

Information provided by summary() depends on the type of data, by column:

carat	**cut**	**color**	**price**
Min. :0.2000	Fair : 1610	D: 6775	Min. : 326
1st Qu.:0.4000	Good : 4906	E: 9797	1st Qu.: 950
Median :0.7000	Very Good:12082	F: 9542	Median : 2401
Mean :0.7979	Premium :13791	G:11292	Mean : 3933
3rd Qu.:1.0400	Ideal :21551	H: 8304	3rd Qu.: 5324
Max. :5.0100		I: 5422	Max. :18823
		J: 2808	

numeric data:
statistical summary

categorical (factor) data:
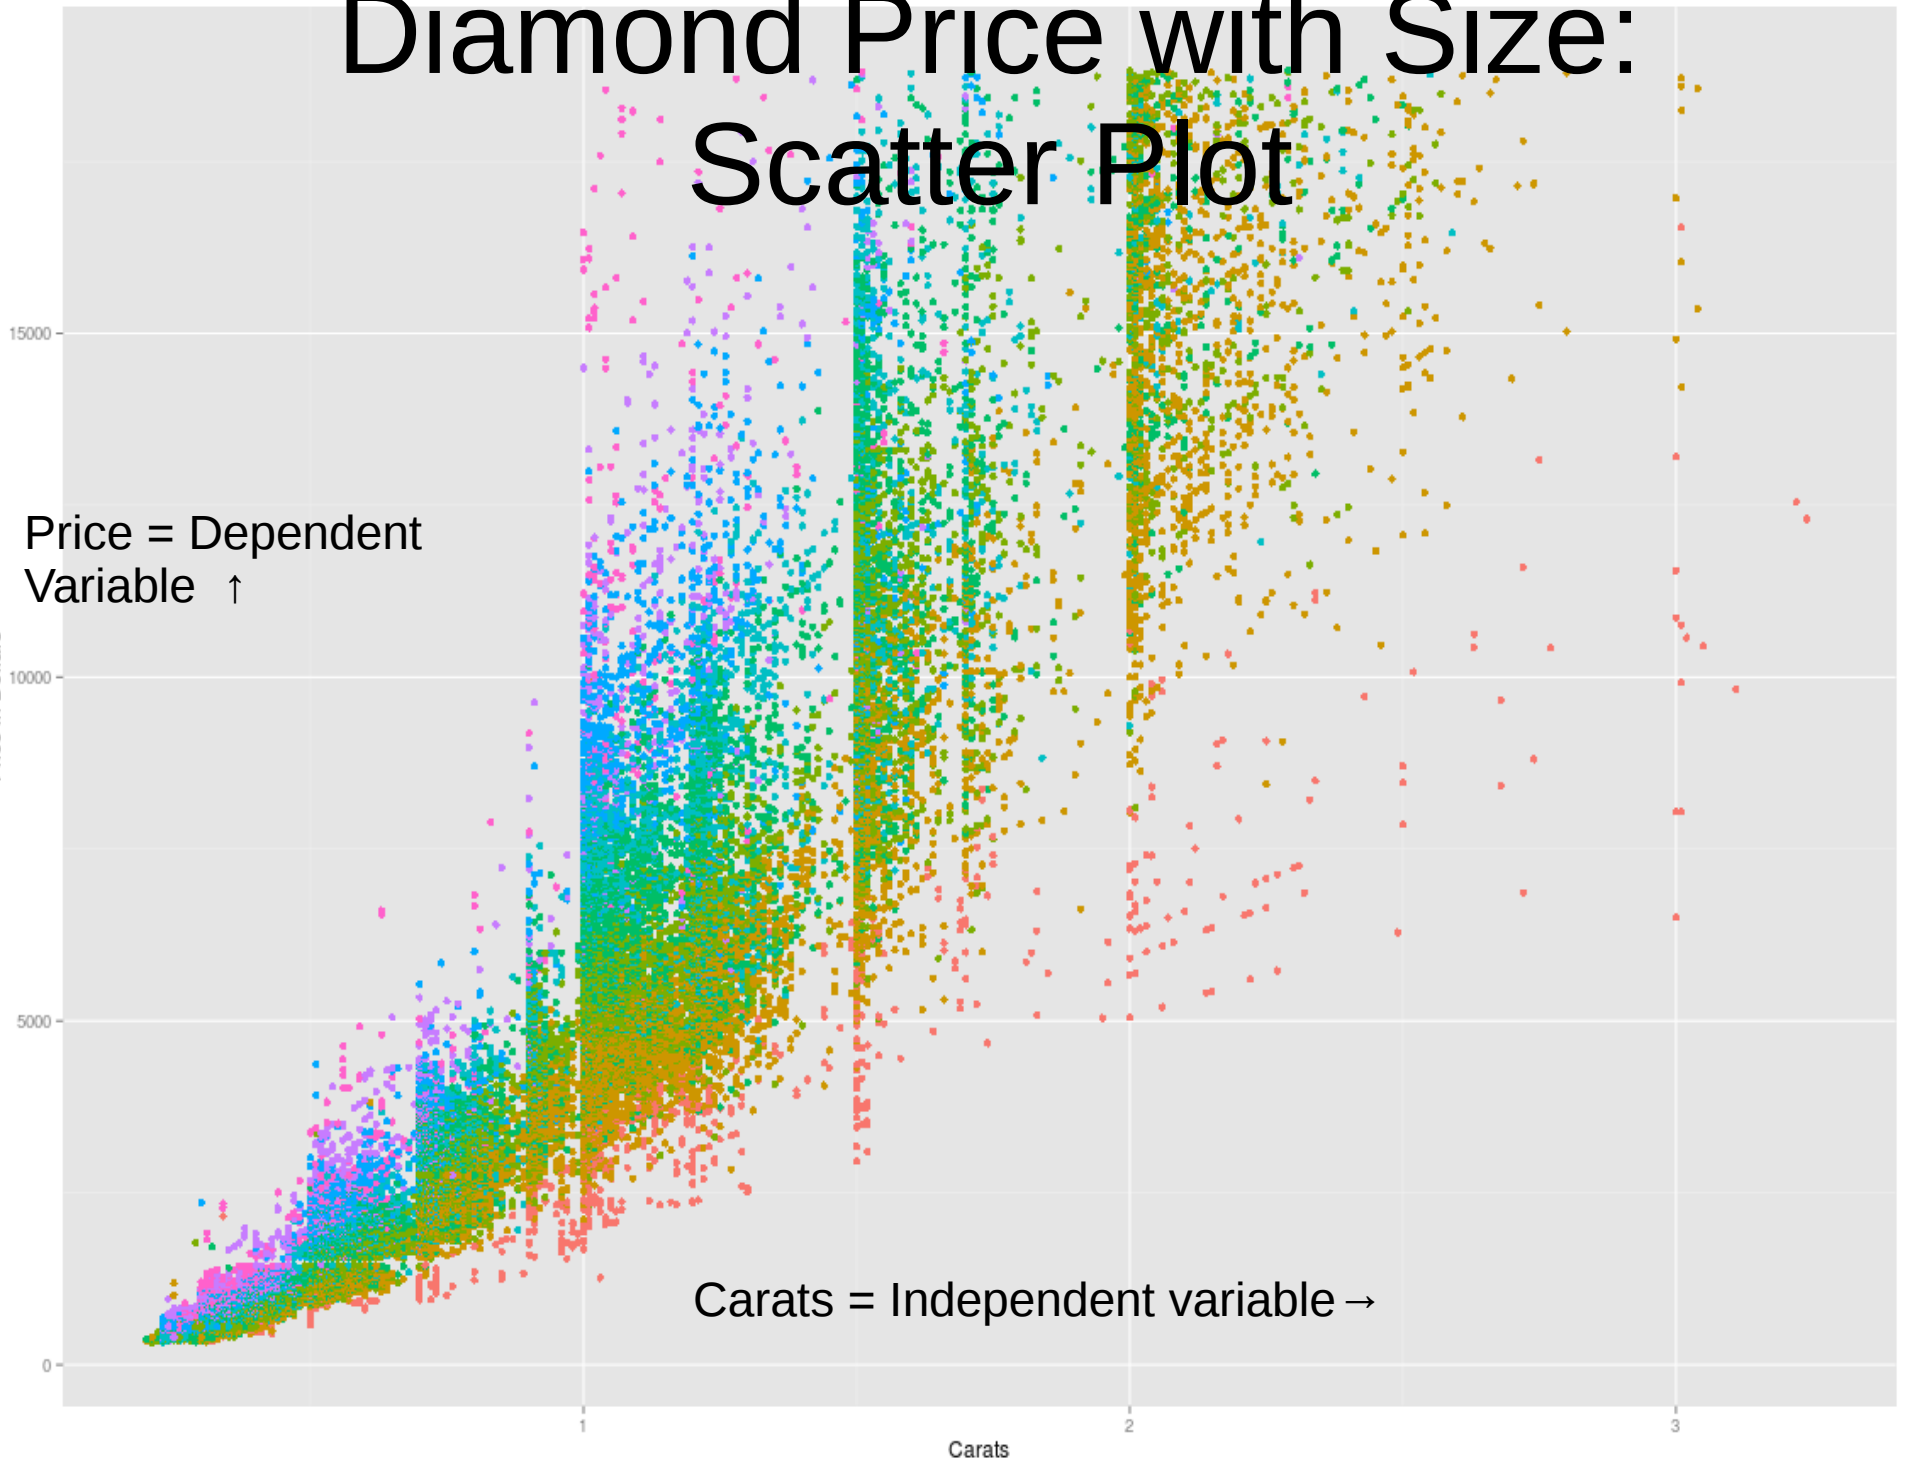counts

Diamond Price with Size: Scatter Plot

Price = Dependent Variable ↑

Price in Dollars

15000

10000

5000

0

Carats = Independent variable →

Carats

1   2   3

clarity

- I1
- SI2
- SI1
- VS2
- VS1
- VVS2
- VVS1
- IF

table() Function

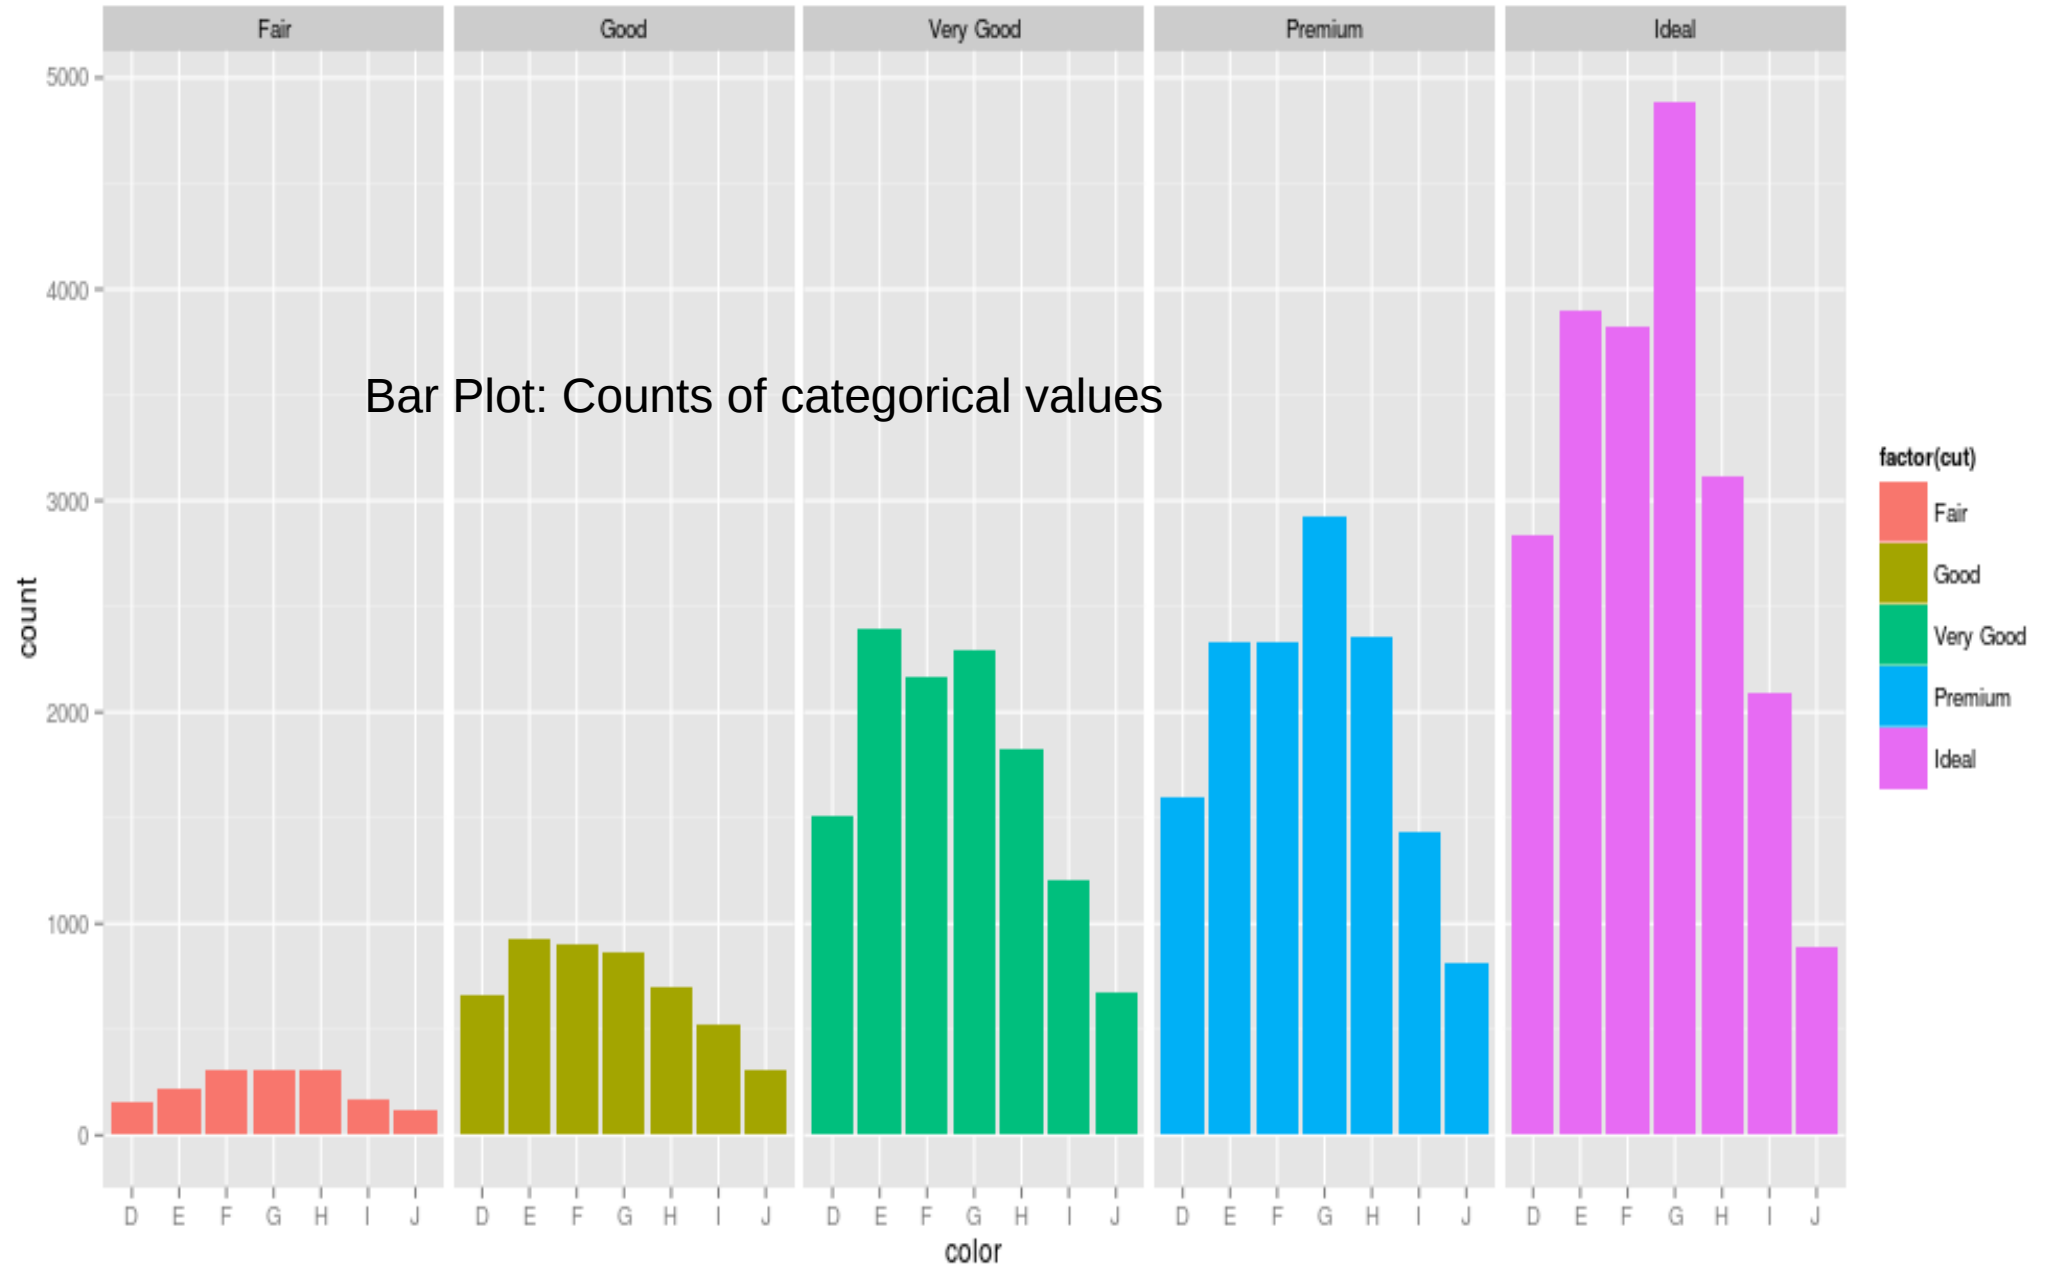Contingency table: counts of categorical values for selected columns

> table(diamonds$cut, diamonds$color)

	D	E	F	G	H	I	J
Fair	163	224	312	314	303	175	119
Good	662	933	909	871	702	522	307
Very Good	1513	2400	2164	2299	1824	1204	678
Premium	1603	2337	2331	2924	2360	1428	808
Ideal	2834	3903	3826	4884	3115	2093	896

Diamond Color and Cut

Bar Plot: Counts of categorical values

Correlation

Do the two quantities X and Y vary together?

- Positively: $0 < \rho < 1$
- Or negatively: $-1 < \rho < 0$

$$\rho_{X,Y} = corr(X,Y) = \frac{cov(X,Y)}{\sigma_X \sigma_Y} = \frac{E[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y}$$

A pairwise, *statistical* relationship between quantities

Correlation

$$\rho_{X,Y} = corr(X,Y) = \frac{cov(X,Y)}{\sigma_X \sigma_Y} = \frac{E[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y}$$

NOTE: Correlation does not imply causation...

Looking for Correlations

diamonds data frame: 50,000 diamonds

- carat: weight of the diamond (0.2–5.01)
- table: width of top of diamond relative to widest point (43–95)
- price: price in US dollars
- x: length in mm (0–10.74)
- y: width in mm (0–58.9)
- z: depth in mm (0–31.8)

cor() Function

Look at pairwise, *statistical* relationships between numeric data:

> cor(diamonds[c(1,6:10)])

	carat	table	price	x	y	z
carat	1.0000000	0.1816175	0.9215913	0.9750942	0.9517222	0.9533874
table	0.1816175	1.0000000	0.1271339	0.1953443	0.1837601	0.1509287
price	0.9215913	0.1271339	1.0000000	0.8844352	0.8654209	0.8612494
x	0.9750942	0.1953443	0.8844352	1.0000000	0.9747015	0.9707718
y	0.9517222	0.1837601	0.8654209	0.9747015	1.0000000	0.9520057
z	0.9533874	0.1509287	0.8612494	0.9707718	0.9520057	1.0000000

-1.0: perfectly anticorrelated

↕

0   : uncorrelated

↕

1.0: perfectly correlated

Student Dataset Example

Now we can write some R to perform some descriptive statistics on our student data:

- Contingency table of school and handedness?

- Correlation between age and height?

- Some descriptive statistics about height (only comparing apples to apples!)

OK...but what's the problem here? Small dataset!

Interlude

Complete descriptive statistics exercises.

Open in the RStudio source editor:

<workshop>/exercises/exercises-descriptive-statistics.R